

CS 461 - Artificial Intelligence

Term Project

2048 with Min-Max Tree Search, MCTS, and DDQN

Group 12

Ceren Akyar, Deniz Mert Dilaverler, Berk Çakar, Elifsena Öz,
İpek Öztaş

Contents

- Introduction
 - What is our goal?
 - Environment
 - Technologies & Libraries
- Algorithms
 - Min-Max Tree Search
 - Monte Carlo Tree Search
 - DDQN
- Conclusion
- Workload Distribution
- References
- Demo

Introduction

- Single-player sliding tile puzzle video **game**
- Use **arrow keys** to move the tiles
- When two tiles with the same number touch, they merge into one
- 4x4 **grid**
- **Tiles:** 512, 1024, 2048
- **Scores:** final game score
- **Game over:** no adjacent tiles have the same value [1]




Figure 1. 2048 Game

What is our goal?

- Implement and test **algorithms**:
 - Min-Max Tree Search
 - Monte Carlo Tree Search
 - Double Deep Q-Network (DDQN) (it is possible to implement DDQN to play 2048 [2])
- Improve the **success rate** (previously it was 7% for achieving 2048 [3])
- **Comparative** analysis
- **Metrics**: Average score, maximum score and highest tile

Environment

- OpenAI Gym library wrapper for 2048
- **Reward** = sum of combined tiles
 $R \Rightarrow 16 + 16 = 32$

- **State Space:** 4x4 matrix

Contains values of the tiles in log2 form

Matrix Values:

0	1	0	0
0	4	7	10
0	5	6	4
3	2	2	1

- **Action Space:** up, down, left, and right



Figure 2. Screenshot of the 2048 environment

Technologies & Libraries

Technologies and libraries used in the project:

- Python
- Numpy
- PyTorch
- Gym
- Matplotlib



Figure 3. Logo of PyTorch

Algorithms

Min-Max Tree Search

- Depth: 3
- Heuristics
 - Corner Heuristic
 - Branch Failure Penalty




Figure 4. Corner heuristic tile weights

```
BRANCH_FAILIURE_PENALTY = -100000
```

Figure 5. Penalty for game termination

Min-Max Tree Search Results

```
Total number of episodes: 10
Minmax tree depth: 3
-----
Average score: 24827.2
Minimum score: 11252
Maximum score: 37724
-----
Tile 512 achieved in 10 / 10 episodes(100.0%)
Tile 1024 achieved in 10 / 10 episodes(100.0%)
Tile 2048 achieved in 6 / 10 episodes(60.0%)
Tile 4096 achieved in 0 / 10 episodes(0.0%)
```

Figure 6. Results for depth 3 Min-Max Search Tree

Min-Max Tree Search Result Plots

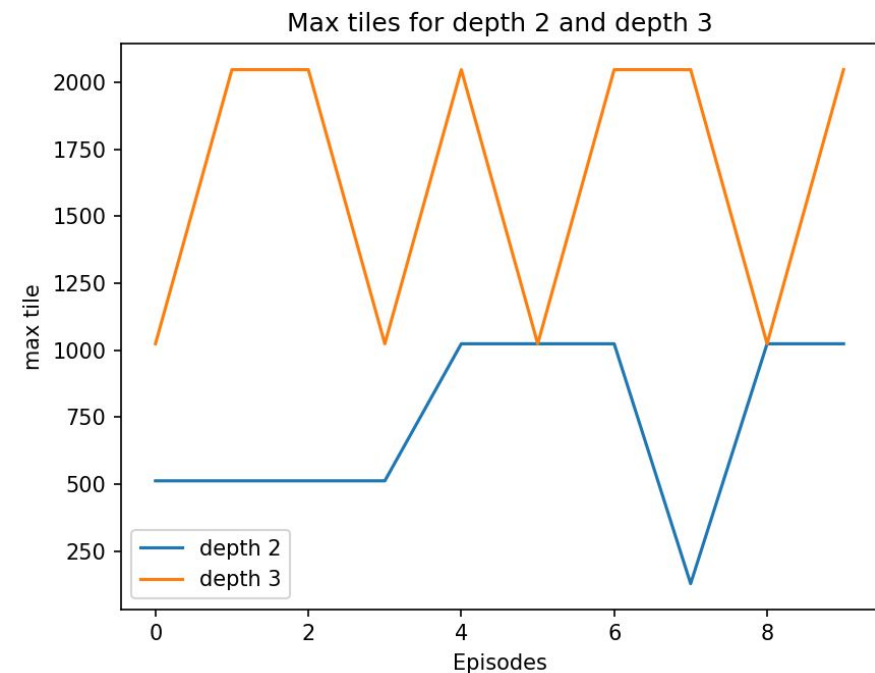


Figure 7. Maximum Tile Reached vs. Number of Episodes




Figure 8. Score Reached vs. Number of Episodes

Min-Max-Tree Search Discussion

- Success rate and highest tile increases with increased depth
 - More foresight
 - Performance problems on higher depths
 - High branching rate despite pruning

Monte Carlo Tree Search Algorithm

- The algorithm was tested for depths 100 and 150
- For depth 100:
 - It can reach the 2048 tile with an 80% success rate
 - It can reach 1024 tile with a 100% success rate
- For depth 150:
 - It can reach the 2048 tile with a 100% success rate
 - It can reach 1024 tile with a 100% success rate

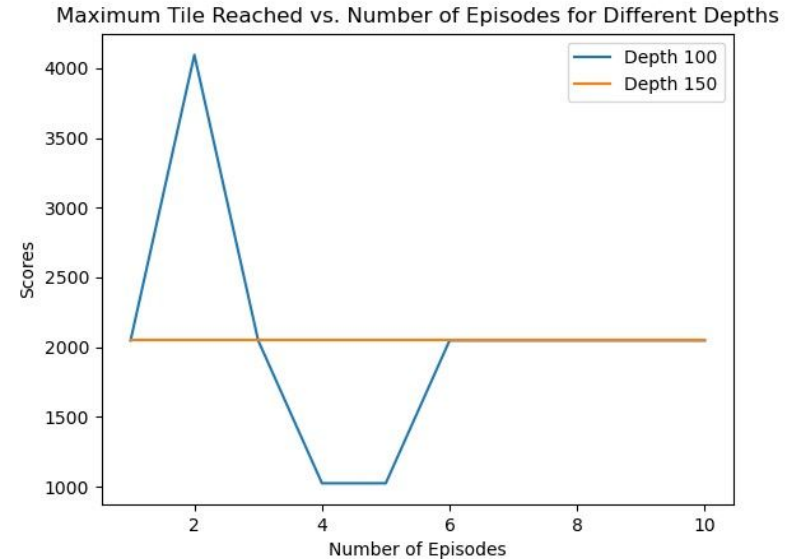


Figure 9. Maximum Tile Reached vs. Number of Episodes

Monte Carlo Tree Search Algorithm

- When the depth increases, the algorithm's overall performance also increases:
 - Because the explored tree depth is increased
 - A more thorough exploration of possible moves and outcomes was allowed

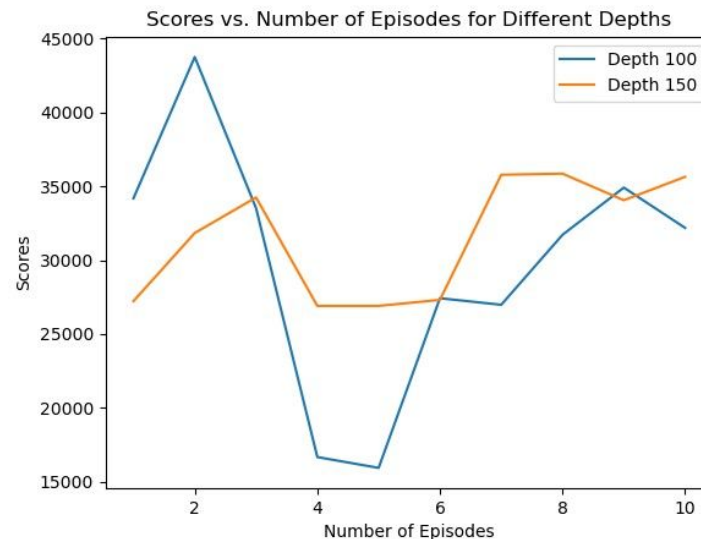


Figure 10. Scores vs. Number of Episodes

Monte Carlo Tree Search Algorithm Comparison




Figure 11. Min-Max Tree Scores

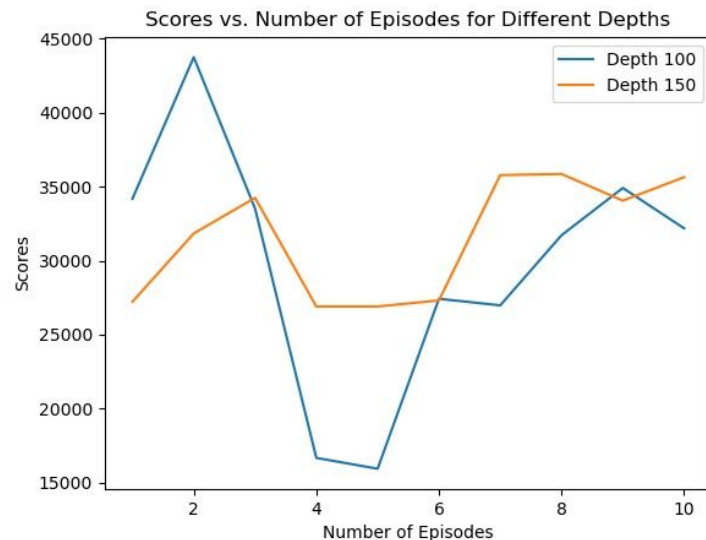


Figure 12. MCTS Scores

- The average performance scores and success rates are higher than the Min-Max tree search algorithm's performance
- Explores promising branches of the search tree
- The Monte Carlo Tree Search algorithm is more suitable for games with uncertainty or stochastic elements

DDQN

- The DDQN Algorithm was tested for:
 - Learning Rates: 0.0001, 0.00025, 0.0005
 - Gamma Values: 0.1, 0.5, 0.8
 - Buffer Sizes: 10000, 50000, 100000
- Results are shown over 1000 episodes (~one million steps) for this algorithm.
- Otherwise, the default values are:
 - LR: 0.0001, Gamma: 0.99, Buffer Size: 10000
 - Epsilon: 1, Epsilon Decay: 0.995, Minimum Epsilon: 0.01

DDQN - Learning Rate Experiments (Max Tiles)




Figure 13. Max Tile vs. Episode (learning rate=1E-4)

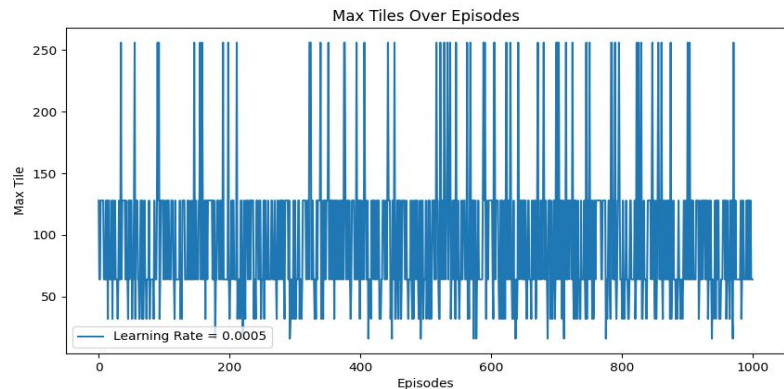


Figure 14. Max Tile vs. Episode (learning rate=5E-4)

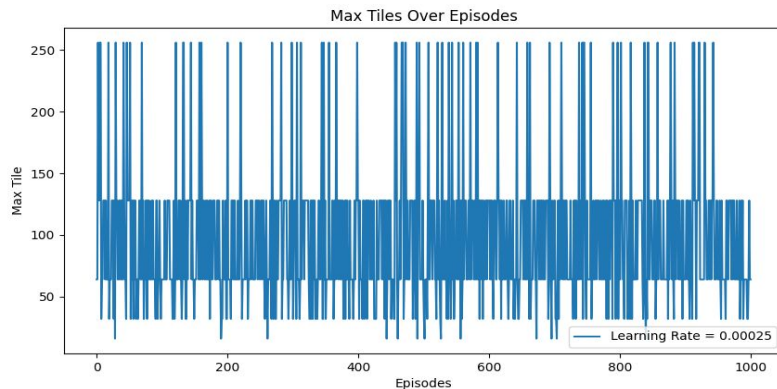


Figure 15. Max Tile vs. Episode (learning rate=2.5E-4)

DDQN - Learning Rate Experiments (Scores)




Figure 16. Score vs. Episode (learning rate=1E-4)




Figure 17. Score vs. Episode (learning rate=5E-4)

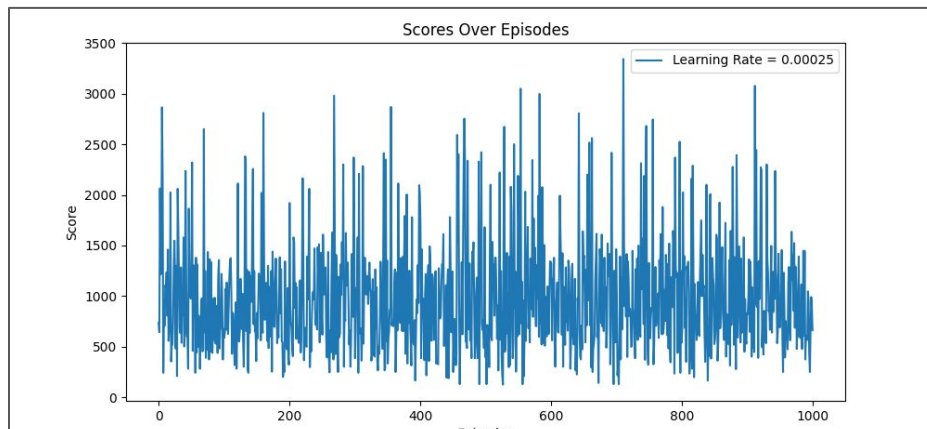


Figure 18. Score vs. Episode (learning rate=2.5E-4)

DDQN - Learning Rate Experiments (Scores)




Figure 19. DDQN Score vs. Episode (with different learning rates)

DDQN - Learning Rate Experiments

Average Scores for Different Learning Rates	
Learning Rate: 0.0001	Average Score: 979.852
Learning Rate: 0.00025	Average Score: 984.764
Learning Rate: 0.0005	Average Score: 958.504
Tiles Achieved for Different Learning Rates	
Learning Rate: 0.0001	
Tile 16 achieved	15 times (1.5%)
Tile 32 achieved	140 times (14.0%)
Tile 64 achieved	378 times (37.8%)
Tile 128 achieved	410 times (41.0%)
Tile 256 achieved	57 times (5.7%)
Learning Rate: 0.00025	
Tile 16 achieved	13 times (1.3%)
Tile 32 achieved	121 times (12.1%)
Tile 64 achieved	418 times (41.8%)
Tile 128 achieved	376 times (37.6%)
Tile 256 achieved	72 times (7.2%)
Learning Rate: 0.0005	
Tile 16 achieved	14 times (1.4%)
Tile 32 achieved	123 times (12.3%)
Tile 64 achieved	416 times (41.6%)
Tile 128 achieved	387 times (38.7%)
Tile 256 achieved	60 times (6.0%)

Figure 20. DDQN Summary Table (with different learning rates)

DDQN - Gamma Value Experiments (Max Tiles)




Figure 21. Max Tile vs. Episode (gamma value=0.1)

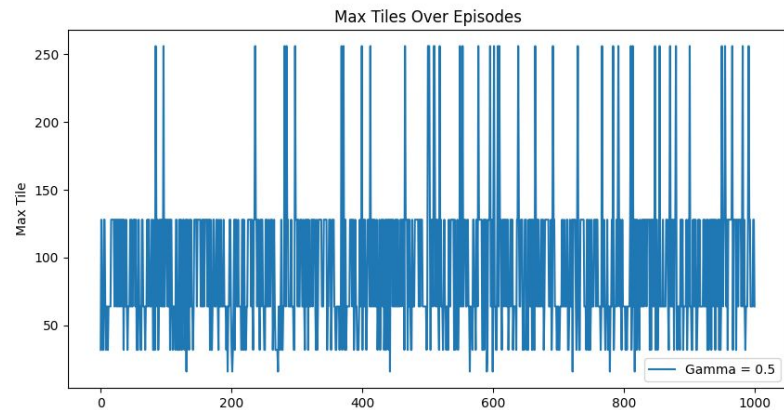


Figure 22. Max Tile vs. Episode (gamma value=0.5)




Figure 23. Max Tile vs. Episode (gamma value=0.8)

DDQN - Gamma Value Experiments (Scores)




Figure 24. Score vs. Episode (gamma value=0.1)




Figure 25. Score vs. Episode (gamma value=0.5)




Figure 26. Score vs. Episode (gamma value=0.8)

DDQN - Gamma Value Experiments (Scores)




Figure 27. DDQN Score vs. Episode (with different gamma values)

DDQN - Gamma Value Experiments

Average Scores for Different Gamma Values	

Gamma: 0.1, Average Score: 928.34	
Gamma: 0.5, Average Score: 924.256	
Gamma: 0.8, Average Score: 963.048	
Tiles Achieved for Different Gamma Values	

Gamma: 0.1	
Tile 16 achieved 26 times (2.6%)	
Tile 32 achieved 146 times (14.6%)	
Tile 64 achieved 408 times (40.8%)	
Tile 128 achieved 363 times (36.3%)	
Tile 256 achieved 57 times (5.7%)	
Gamma: 0.5	
Tile 16 achieved 11 times (1.1%)	
Tile 32 achieved 157 times (15.7%)	
Tile 64 achieved 386 times (38.6%)	
Tile 128 achieved 403 times (40.3%)	
Tile 256 achieved 43 times (4.3%)	
Gamma: 0.8	
Tile 16 achieved 15 times (1.5%)	
Tile 32 achieved 146 times (14.6%)	
Tile 64 achieved 389 times (38.9%)	
Tile 128 achieved 387 times (38.7%)	
Tile 256 achieved 63 times (6.3%)	

Figure 28. DDQN Summary Table (with different gamma values)

DDQN - Buffer Size Experiments (Max Tiles)




Figure 29. Max Tile vs. Episode (batch size=1E4)




Figure 30. Max Tile vs. Episode (batch size=5E4)




Figure 31. Max Tile vs. Episode (batch size=1E5)

DDQN - Buffer Size Experiments (Scores)




Figure 32. Score vs. Episode (batch size=1E4)




Figure 33. Score vs. Episode (batch size=5E4)




Figure 34. Score vs. Episode (batch size=1E5)

DDQN - Buffer Size Experiments (Scores)




Figure 35. DDQN Score vs. Episode (with different batch sizes)

DDQN - Buffer Size Experiments

Average Scores for Different Buffer Sizes	
Buffer Size: 10000	Average Score: 979.852
Buffer Size: 50000	Average Score: 989.64
Buffer Size: 100000	Average Score: 976.432
Tiles Achieved for Different Buffer Sizes	
Buffer Size: 10000	
Tile 16	achieved 15 times (1.5%)
Tile 32	achieved 140 times (14.0%)
Tile 64	achieved 378 times (37.8%)
Tile 128	achieved 410 times (41.0%)
Tile 256	achieved 57 times (5.7%)
Buffer Size: 50000	
Tile 16	achieved 7 times (0.7%)
Tile 32	achieved 129 times (12.9%)
Tile 64	achieved 390 times (39.0%)
Tile 128	achieved 407 times (40.7%)
Tile 256	achieved 67 times (6.7%)
Buffer Size: 100000	
Tile 8	achieved 1 times (0.1%)
Tile 16	achieved 12 times (1.2%)
Tile 32	achieved 125 times (12.5%)
Tile 64	achieved 406 times (40.6%)
Tile 128	achieved 401 times (40.1%)
Tile 256	achieved 55 times (5.5%)

Figure 36. DDQN Summary Table (with different gamma values)

DDQN - Discussion

- DDQN performs worse than both Min-Max Tree Search and MCTS. There might be several underlying reasons for that:
 - Complexity of the Game:
 - 2048 involves a significant amount of strategic depth.
 - MCTS and Min-Max Tree Search may be performing due to their lookahead capabilities.
 - DDQN, being a value-based method, might fail to plan strategically. It is much easier for DDQN to learn to play control/instant action games (e.g. Pong)

DDQN - Discussion

- Other threats for the DDQN's performance can include:
 - Neural Network Architecture
 - Hyperparameters
 - Training Duration
 - Rewarding System

Conclusion

- We explored the ways to “solve” the game 2048 with AI algorithms.
- We used Min-Max Tree Search, MCTS, and DDQN algorithms
- Performances: MCTS > Min-Max Tree Search > DDQN
 - Due to the previously discussed points, plus the fact that:
 - Tree-based algorithms handle the future planning better, and perform well on strategical games. Hence, 2048 is not an exception.

Workload Distribution

- Ceren Akyar: MCTS experiments, plotting & inferring the results
- Deniz Mert Dilaverler: Implementation of Min-Max Tree Search and its experiments
- Berk Çakar: Implementation of MCTS, DDQN algorithms
- Elifsena Öz: DDQN experiments, plotting & inferring the results
- İpek Öztaş: DDQN experiments, plotting & inferring the results

References

- [1] “2048 (video game),” Wikipedia. Oct. 20, 2023. Accessed: Nov. 26, 2023. [Online]. Available: [https://en.wikipedia.org/w/index.php?title=2048_\(video_game\)&oldid=1185818281](https://en.wikipedia.org/w/index.php?title=2048_(video_game)&oldid=1185818281)
- [2] L. H. Chan, “Playing 2048 with Deep Q-Learning (with pytorch implementation),” Medium, Oct. 20, 2023. [Online]. Available: <https://medium.com/@qwert12500/playing-2048-with-deep-q-learning-with-pytorch-implementation-4313291efe61>. [Accessed: Oct. 20, 2023].
- [3] A. Goga, *Reinforcement learning in 2048 game*, Oct. 2017. Accessed: Oct. 20, 2023. [Online]. Available: <http://cogsci.fmph.uniba.sk/~farkas/theses/adrian.goga.bak18.pdf>

Demo Video

score: 20

best: 20

max tile: 8

episode: 1

2048 count: 0

2	4		
8		2	
8			
2			

score: 180

best: 180

max tile: 32

episode: 1

2048 count: 0

8	16	8	
32	2		
		2	

score: 180

best: 180

max tile: 32

episode: 1

2048 count: 0

8	16	8	
32	2		
		2	